Incident Response and On-Call Practice

What Counts As An Incident

An incident is any unplanned disruption that affects users or threatens to. The definition is deliberately broad, because the cost of declaring an incident that turns out to be minor is far lower than the cost of treating a real outage as routine work for several hours.

Severity should be assessed by user impact rather than by technical interest. A subtle bug affecting all users outranks a dramatic failure in a system nobody is currently depending on.

Roles During An Incident

The incident commander coordinates and decides; they do not debug. The most common failure in incident response is the most knowledgeable person taking command and then disappearing into a terminal, leaving nobody tracking the overall picture.

A separate communications role keeps stakeholders informed on a predictable cadence. Without one, responders are interrupted continuously for status updates, which lengthens the outage.

Mitigate Before Diagnosing

Restoring service takes priority over understanding the cause. Rolling back a recent deployment, shifting traffic, or disabling a feature flag are all acceptable even when it is not yet known whether they address the true cause.

Diagnosis performed while users are affected is diagnosis performed under time pressure, which is when mistakes are made. Preserving logs and metrics for later analysis costs little and protects the investigation.

Writing A Blameless Postmortem

A postmortem assumes that everyone acted reasonably given the information they had at the time. The purpose is to find the conditions that made the failure possible, because those conditions will still be present after the individual has been more careful.

Naming an individual as the cause reliably stops the flow of information. Once engineers expect to be blamed, near-misses stop being reported, and the organisation loses its cheapest source of warning.

Action items should be specific and owned. An item reading improve monitoring will not be done; an item naming a specific alert, a threshold and an owner might be.

Effective Alerting Practice

Every alert should be actionable and urgent. An alert that fires regularly and is routinely ignored is worse than no alert, because it trains responders to dismiss the whole channel.

Alert on symptoms that users experience rather than on internal causes. High processor usage may be entirely normal, while a rise in request latency is a problem regardless of which internal cause produced it.